

# Word2Vec-以 gensim 訓練中文詞向量

## 參考及引用資料來源

- [1] [zake7749-使用 gensim 訓練中文詞向量 \(http://zake7749.github.io/2016/08/28/word2vec-with-gensim/\)](http://zake7749.github.io/2016/08/28/word2vec-with-gensim/)
- [2] [gensim/corpora/wikicorpus \(https://radimrehurek.com/gensim/corpora/wikicorpus.html\)](https://radimrehurek.com/gensim/corpora/wikicorpus.html)
- [Word2Vec的簡易教學與參數調整指南 \(https://www.kaggle.com/jerrykuo7727/word2vec\)](https://www.kaggle.com/jerrykuo7727/word2vec)
- [zhconv \(https://pypi.org/project/zhconv/\)](https://pypi.org/project/zhconv/)
- [jieba \(https://pypi.org/project/jieba/\)](https://pypi.org/project/jieba/)

In [2]:

```
%load_ext memory_profiler
!pip install -q zhconv
```

確認相關 Packages

In [3]:

```
import os

# Packages
import gensim
import jieba
import zhconv
from gensim.corpora import WikiCorpus
from datetime import datetime as dt
from typing import List

if not os.path.isfile('dict.txt.big'):
    !wget https://github.com/fxsjy/jieba/raw/master/extra_dict/dict.txt.big
    jieba.set_dictionary('dict.txt.big')

print("gensim", gensim.__version__)
print("jieba", jieba.__version__)
```

```
gensim 4.1.2
jieba 0.42.1
```

## 準備中文訓練文本

訓練文本來源: [維基百科資料庫](#)

<https://zh.wikipedia.org/wiki/Wikipedia:%E6%95%B0%E6%8D%A>

要訓練詞向量，第一步當然是取得資料集。由於 word2vec 是基於非監督式學習，訓練集一定一定要越大越好，語料涵蓋的越全面，訓練出來的結果也會越漂亮。[1]  
(<http://zake7749.github.io/2016/08/28/word2vec-with-gensim/>)

- [zhwiki-20210101-pages-articles.xml.bz2](https://dumps.wikimedia.org/zhwiki/20210101/zhwiki-20210101-pages-articles.xml.bz2) (<https://dumps.wikimedia.org/zhwiki/20210101/zhwiki-20210101-pages-articles.xml.bz2>) (1.9 GB)

目前已經使用另一份 Notebook ([維基百科中文語料庫 zhWiki 20210101](https://www.kaggle.com/bbqlp33/zhwiki-20210101)  
(<https://www.kaggle.com/bbqlp33/zhwiki-20210101>)) 下載好中文維基百科語料，並可以直接引用

In [17]:

```
#ZhWiki = "/kaggle/input/zhwiki-20210101/zhwiki-20210101-pages-articles.xml.bz2"
ZhWiki = r"C:\Users\user\Desktop\exercise\NLP\zhwiki-20230501-pages-articles.xml.bz2"
#"C:\Users\user\Desktop\exercise\NLP\zhwiki-20230501-pages-articles.xml.bz2"
!du -sh $ZhWiki
!md5sum $ZhWiki
!file $ZhWiki
```

'du' 不是內部或外部命令、可執行的程式或批次檔。  
'md5sum' 不是內部或外部命令、可執行的程式或批次檔。  
'file' 不是內部或外部命令、可執行的程式或批次檔。

## 中文文本前處理

在正式訓練 Word2Vec 之前，其實涉及了文本的前處理，本篇的處理包括如下三點 (而實務上對應的不同使用情境，可能會有不同的前處理流程):

- 簡轉繁: [zhconv](https://pypi.org/project/zhconv/) (<https://pypi.org/project/zhconv/>)
- 中文斷詞: [jieba](https://pypi.org/project/jieba/) (<https://pypi.org/project/jieba/>)
- 停用詞

## 簡繁轉換

wiki 文本其實摻雜了簡體與繁體中文，比如「数学」與「數學」，這會被 word2vec 當成兩個不同的詞。[1]  
(<http://zake7749.github.io/2016/08/28/word2vec-with-gensim/>)

所以我們在斷詞前，需要加上簡繁轉換的手續

以下範例使用了較輕量的 Package [zhconv](https://pypi.org/project/zhconv/) (<https://pypi.org/project/zhconv/>)，  
若需要更高的精準度，則可以參考 [OpenCC](https://github.com/BYVoid/OpenCC) (<https://github.com/BYVoid/OpenCC>)

In [6]:

```
zhconv.convert("这原本是一段简体中文", "zh-tw")
```

Out[6]:

```
'這原本是一段簡體中文'
```

## 中文斷詞

使用 [jieba](https://pypi.org/project/jieba/) (<https://pypi.org/project/jieba/>) `jieba.cut` 來進行中文斷詞，並簡單介紹 `jieba` 的兩種分詞模式：

- `cut_all=False` **精確模式**，試圖將句子最精確地切開，適合文本分析；
- `cut_all=True` **全模式**，把句子中所有的可以成詞的詞語都掃描出來，速度非常快，但是不能解決歧義；

而本篇文本訓練採用**精確模式** `cut_all=False`

In [7]:

```
seg_list = jieba.cut("我来到北京清华大学", cut_all=True)
print("Full Mode: " + "/ ".join(seg_list))  # 全模式

seg_list = jieba.cut("我来到北京清华大学", cut_all=False)
print("Default Mode: " + "/ ".join(seg_list))  # 精確模式
```

```
Building prefix dict from C:\Users\user\Desktop\exercise\NLP\dict.txt.big
...
Dumping model to file cache C:\Users\user\AppData\Local\Temp\jieba.u1f62a
42fe1738d51f5bb243f550ba545.cache
Loading model cost 1.276 seconds.
Prefix dict has been built successfully.
```

```
Full Mode: 我/ 来到/ 北京/ 清华/ 清华大学/ 华大/ 大学
Default Mode: 我/ 来到/ 北京/ 清华大学
```

In [8]:

```
print(list(jieba.cut("中英夾雜的example·Word2Vec應該很interesting吧?")))
```

```
['中', '英', '夾雜', '的', 'example', '·', 'Word2Vec', '應該', '很', 'inte', 'resting', '吧', '?']
```

## 引入停用詞表

停用詞就是像英文中的 **the,a,this**，中文的**你我他**，與其他詞相比顯得不怎麼重要，對文章主題也無關緊要的，

是否要使用停用詞表，其實還是要看你的應用，也有可能保留這些停用詞更能達到你的目標。[1]

(<http://zake7749.github.io/2016/08/28/word2vec-with-gensim/>)

- [Is it compulsory to remove stop words with word2vec?](https://www.quora.com/Is-it-compulsory-to-remove-stop-words-with-word2vec) (<https://www.quora.com/Is-it-compulsory-to-remove-stop-words-with-word2vec>)
- [The Effect of Stopword Filtering prior to Word Embedding Training](https://stats.stackexchange.com/questions/201372/the-effect-of-stopword-filtering-prior-to-word-embedding-training) ([https://stats.stackexchange.com/questions/201372/the-effect-of-stopword-filtering-prior-to-word-](https://stats.stackexchange.com/questions/201372/the-effect-of-stopword-filtering-prior-to-word-embedding-training)

[embedding-training](#))

以下範例還是示範引入停用詞表，而停用詞表網路上有各種各樣的資源

剛好 kaggle 環境預設有裝 [spacy](https://pypi.org/project/spacy/) (<https://pypi.org/project/spacy/>)，

就順道引用 spacy 提供的停用詞表吧 (實務上 stopwords 應為另外準備好且檢視過的靜態文檔)

In [10]:

```
import spacy

# 下載語言模組
spacy.cli.download("zh_core_web_sm") # 下載 spacy 中文模組
spacy.cli.download("en_core_web_sm") # 下載 spacy 英文模組

nlp_zh = spacy.load("zh_core_web_sm") # 載入 spacy 中文模組
nlp_en = spacy.load("en_core_web_sm") # 載入 spacy 英文模組

# 印出前20個停用詞
print('--\n')
print(f"中文停用詞 Total={len(nlp_zh.Defaults.stop_words)}: {list(nlp_zh.Defaults.stop_words)}")
print("--")
print(f"英文停用詞 Total={len(nlp_en.Defaults.stop_words)}: {list(nlp_en.Defaults.stop_words)}")
```

✓ Download and installation successful

You can now load the package via spacy.load('zh\_core\_web\_sm')

✓ Download and installation successful

You can now load the package via spacy.load('en\_core\_web\_sm')

--

中文停用詞 Total=1891: ['诚然', '同', '怎样', '■', '变成', '呜呼', '一般', '犹且', '是以', '论说', '什么', '如次', '倘若', '∧', '咳', '至今', '多么', '最后', '啊哟', '且不说'] ...

--

英文停用詞 Total=326: ['re', 'empty', 'forty', 'ten', 'few', 'thru', 's', 'meanwhile', 'no', 'perhaps', 'upon', 'although', 'call', 'it', 'm', 'off', 'same', 'sometimes', 'eight', 'using'] ...

In [11]:

```
STOPWORDS = nlp_zh.Defaults.stop_words | \
            nlp_en.Defaults.stop_words | \
            set(["\n", "\r\n", "\t", " ", ""])
print(len(STOPWORDS))

# 將簡體停用詞轉成繁體，擴充停用詞表
for word in STOPWORDS.copy():
    STOPWORDS.add(zhconv.convert(word, "zh-tw"))

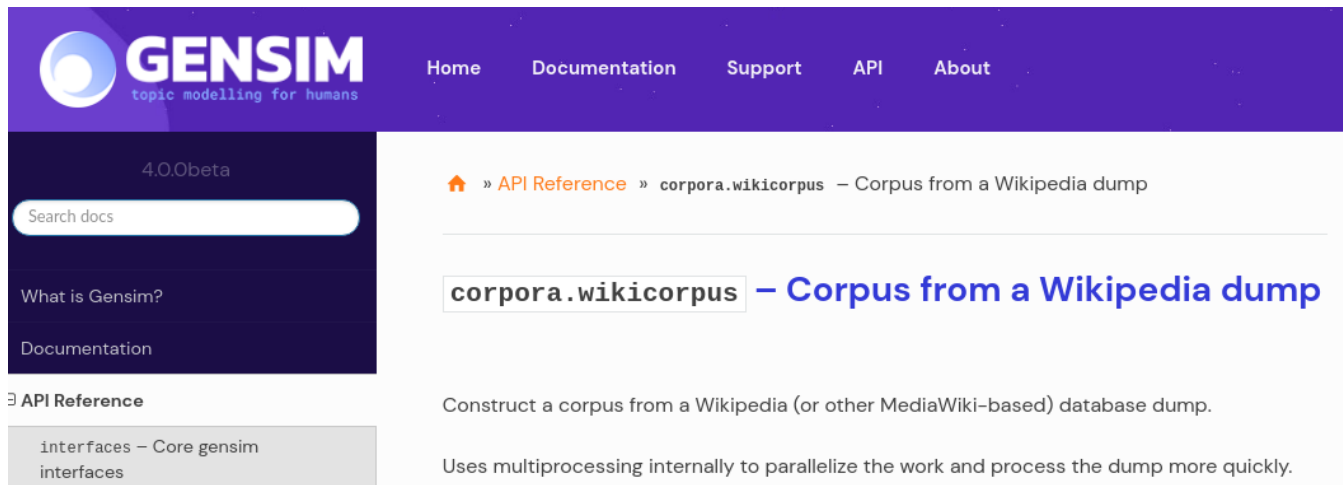
print(len(STOPWORDS))
```

2222

3005

## 讀取 wiki 語料庫，並且進行前處理和斷詞

維基百科 ( [wiki.xml.bz2](#) ) 下載好後，先別急著解壓縮，因為這是一份 xml 文件，裏頭佈滿了各式各樣的標籤，我們得先想辦法送走這群不速之客，不過也別太擔心，gensim 早已看穿了一切，藉由調用 [wikiCorpus](#) (<https://radimrehurek.com/gensim/corpora/wikicorpus.html>)，我們能很輕鬆的只取出文章的標題和內容。[1] (<http://zake7749.github.io/2016/08/28/word2vec-with-gensim/>)



[2] (<https://radimrehurek.com/gensim/corpora/wikicorpus.html>).

Supported dump formats:

- <LANG>wiki-<YYYYMMDD>-pages-articles.xml.bz2
- <LANG>wiki-latest-pages-articles.xml.bz2

In [12]:

```
def preprocess_and_tokenize(
    text: str, token_min_len: int=1, token_max_len: int=15, lower: bool=True) -> List[str]:
    if lower:
        text = text.lower()
    text = zhconv.convert(text, "zh-tw")
    return [
        token for token in jieba.cut(text, cut_all=False)
        if token_min_len <= len(token) <= token_max_len and \
           token not in STOPWORDS
    ]
```

In [13]:

```
print(preprocess_and_tokenize("歐幾里得，西元前三世紀的古希臘數學家，現在被認為是幾何之父，此畫拉斐爾"))
print(preprocess_and_tokenize("我来到北京清华大学"))
print(preprocess_and_tokenize("中英夾雜的example，Word2Vec應該很interesting吧?"))
```

```
['歐幾', '裡得', '西元前', '世紀', '古希臘', '數學家', '幾何', '父', '此畫', '拉斐爾']
['來到', '北京', '清華大學']
['中', '英', '夾雜', 'example', 'word2vec', 'interesting']
```

In [ ]:

```
%%time
%%memit

print(f"Parsing {ZhWiki}...")
wiki_corpus = WikiCorpus(ZhWiki, tokenizer_func=preprocess_and_tokenize, token_min_len=1
```

In [ ]:

In [\*]:

```
g = wiki_corpus.get_texts()
print(next(g)[:10])
print(next(g)[:10])
print(next(g)[:10])

# print(jieba.Lcut("".join(next(g))[:50]))
# print(jieba.Lcut("".join(next(g))[:50]))
```

初始化 `WikiCorpus` 後，能藉由 `get_texts()` 可迭代每一篇文章，它所回傳的是一個 `tokens list`，我以空白符將這些 `tokens` 串接起來，統一輸出到同一份文字檔裡。這邊要注意一件事，`get_texts()` 受 `article_min_tokens` 參數的限制，只會回傳內容長度大於 **50** (default) 的文章。

- **article\_min\_tokens** (*int, optional*) – Minimum tokens in article. Article will be ignored if number of tokens is less.

秀出前 3 篇文章的前10 個 token

## 將處理完的語料集存下來，供後續使用

In [\*]:

```
WIKI_SEG_TXT = "wiki_seg.txt"

generator = wiki_corpus.get_texts()

with open(WIKI_SEG_TXT, "w", encoding='utf-8') as output:
    for texts_num, tokens in enumerate(generator):
        output.write(" ".join(tokens) + "\n")

        if (texts_num + 1) % 100000 == 0:
            print(f"[{str(dt.now()):.19}] 已寫入 {texts_num} 篇斷詞文章")
```

## 訓練 Word2Vec

In [14]:

```
%%time

from gensim.models import word2vec
import multiprocessing

max_cpu_counts = multiprocessing.cpu_count()
word_dim_size = 300 # 設定 word vector 維度
print(f"Use {max_cpu_counts} workers to train Word2Vec (dim={word_dim_size})")

# 讀取訓練語句
sentences = word2vec.LineSentence(WIKI_SEG_TXT)

# 訓練模型
model = word2vec.Word2Vec(sentences, size=word_dim_size, workers=max_cpu_counts)

# 儲存模型
output_model = f"word2vec.zh.{word_dim_size}.model"
model.save(output_model)
```

```
Use 4 workers to train Word2Vec (dim=300)
CPU times: user 3h 8min 9s, sys: 1min 25s, total: 3h 9min 34s
Wall time: 1h 4min 34s
```

儲存的模型總共會產生三份檔案

In [15]:

```
! ls word2vec.zh*
```

```
word2vec.zh.300.model
word2vec.zh.300.model.trainables.syn1neg.npy
word2vec.zh.300.model.wv.vectors.npy
```

In [16]:

```
!du -sh word2vec.zh*
```

```
71M      word2vec.zh.300.model
1.3G     word2vec.zh.300.model.trainables.syn1neg.npy
1.3G     word2vec.zh.300.model.wv.vectors.npy
```

## 查看模型以及詞向量實驗

模型其實就是巨大的 Embedding Matrix

In [17]:

```
print(model.wv.vectors.shape)
model.wv.vectors
```

(1138562, 300)

Out[17]:

```
array([[ -1.9005029e+00, -2.8552267e-01, -1.6861261e+00, ...,
        -2.3989561e+00,  7.9282230e-01, -1.7689761e+00],
       [-8.9374363e-01,  2.7187495e+00, -2.8586307e+00, ...,
        7.6278639e-01, -1.2364342e+00,  4.6035990e-01],
       [-1.3627478e+00,  9.0678096e-01, -2.4168897e+00, ...,
        1.9464742e-01,  2.3467582e-01, -4.3651393e-01],
       ...,
       [ 1.1225304e-01, -2.3689542e-02,  5.3460799e-02, ...,
        1.8911289e-02,  6.1728880e-02, -2.9928887e-02],
       [ 1.0599125e-01,  1.4494479e-03, -7.2507304e-03, ...,
        3.7081163e-02,  1.5778864e-02, -1.6071750e-02],
       [-5.2874412e-02, -6.3391119e-02, -1.4798551e-02, ...,
        -7.8865103e-02,  3.2674141e-02, -2.9911553e-02]], dtype=float32)
```

收錄的詞彙

In [18]:

```
print(f"總共收錄了 {len(model.wv.vocab)} 個詞彙")

print("印出 20 個收錄詞彙:")
print(list(model.wv.vocab.keys())[:10])
```

總共收錄了 1138562 個詞彙

印出 20 個收錄詞彙:

['歐幾里得', '西元前', '世紀', '古希臘', '數學家', '認為', '幾何', '父', '此  
畫', '為']

詞彙的向量



In [19]:

```
vec = model.wv['數學家']  
print(vec.shape)  
vec
```

(300,)

Out[19]:

沒見過的詞彙

In [20]:

```
word = "這肯定沒見過 "  
  
# 若強行取值會報錯  
try:  
    vec = model.wv[word]  
except KeyError as e:  
    print(e)
```

"word '這肯定沒見過 ' not in vocabulary"

## 查看前 10 名相似詞

model.wv.most\_similar 的 topn 預設為 10

In [21]:

```
model.wv.most_similar("飲料", topn=10)
```

Out[21]:

```
[('飲品', 0.7986035346984863),  
 ('果汁', 0.6965157389640808),  
 ('含酒精', 0.6784194707870483),  
 ('提神', 0.6489840149879456),  
 ('酒精類', 0.6488814353942871),  
 ('罐裝', 0.6478314995765686),  
 ('軟飲料', 0.6374701261520386),  
 ('啤酒', 0.6282455921173096),  
 ('酒類', 0.6256542801856995),  
 ('優格', 0.6210787892341614)]
```

```
array([-0.15254694, -1.7895197, 2.8389933, 0.30419526, 0.94147015,
       -0.924163, 0.05298484, -0.13174261, -1.4274364, -0.85435706,
       0.04792711, 0.0373514, 0.05875092, 1.6739838, 2.3618689,
       2.0175905, -1.5817889, -0.8472086, 1.1092311, -0.24629189,
```

```
Out[22]: 0.49799377, 0.36100635, -0.281036, -0.5806378, -1.8780046,
         0.2837012, -0.11682022, -0.22787744, 0.11489099, -0.14591609,
         [('truck', 0.20460258549324025645), -0.01348399, -1.0098934, 3.2184417,
          ('motorcycle', 0.70465193974071076163), 1.9969889, 0.14035068, 3.6188982,
          ('motorcycle', 0.70465193974071076163), 1.9969889, 0.14035068, 3.6188982,
          ('seat', 0.0864599225860147054804, 1.7810897, 1.9077462, 0.33889568,
          ('vehicle', 0.4860808390409040441065), -1.7290001, 0.52984065, 1.971243,
          ('chevrolet', 0.1042730.6203585067060956), 0.15407254, 1.0061789, 0.36432067,
          ('wagon', 0.1096308834105709218088, -0.6530562, 0.5096923, 1.6476531,
          ('volkswagen', 0.283320.618094510870604795), 0.710713, -2.182122, -1.2232319,
          ('cab', 0.0161604086560796840824, 0.3732413, 0.38064095, 3.7629342,
          ('carriage', 0.5054040120896788406504), 0.3355076, -0.82843256, -0.01899206,
          1.3599714, -0.3263021, 0.5632072, -1.6275458, 0.18896504,
```

```
In [23]: 2.8486328, 0.4079449, 0.76394415, 0.39884079, -1.1291069,
         1.616785, -2.2050192, -0.68297696, -0.7727057, 1.8106886,
         model.wv.most_similar("4100050k") 0.9511317, 0.3185586, -2.787385,
```

```
Out[23]: -2.046976, 1.1417929, 0.65914726, -0.09776347, 1.291978,
         0.9779469, -2.1436791, 1.177742, 2.2546558, -1.7125031,
         -0.95545936, -0.11377065, -1.223585, -1.3839793, -1.6620296,
         [('臉書', 0.0921058917212028532), 2.6440482, -0.19194634, 0.38509196,
          ('專頁', 0.1746109224557817689), 0.23664996, 1.3536583, 2.3723216,
          ('面書', 0.0491000787734982530), -1.3566401, -1.283516, 0.4361322,
          ('instagram', 0.083003731069080821477), 2.1489894, -0.470945, 0.28165355,
          ('貼文', 0.065046813729667662), -0.47257465, -0.8030565, -1.4145786,
          ('推特', 0.055827963647840886), -3.5372381, 1.4158623, 0.3096819,
          ('粉絲團', 0.7040821428537368474), 1.6208903, -1.2401619, -0.13551526,
          ('twitter', 0.0001056033058348449), 0.4481448, -0.4172336, -1.206034,
          ('網誌', 0.235708784349415728), -1.966441, -0.44706354, 0.34252086,
          ('youtube', 0.645046867870430208142), -0.6640831, -0.80406773, -1.0047605,
```

```
In [24]: 0.25708237, -1.2658246, 0.72502005, -0.16557129, -2.1745164,
         -0.37656462, -1.1152672, 1.0244526, 1.1565366, -1.3018371,
         -0.61735666, -0.47080624, 1.0154963, 0.27704978, -0.43282175,
         model.wv.most_similar("01108066", 0.9108066, -1.0534698, 0.8483542, -2.0334537,
```

```
Out[24]: -1.710315, -1.3449622, 2.2180285, -0.22329706, 0.3695973,
         1.5329542, -0.46323976, -1.4827193, 0.73403674, -0.34677625,
         1.0670941, 0.88510114, -1.906598, -0.5413286, -3.2423143,
         [('欺詐', 0.0974238093640365609), -2.0051272, 0.8977296, -0.39034557,
          ('詐騙', 0.033067189198289868), 0.86995924, -0.42781776, -0.79106665,
          ('竊盜', 0.050685684191513062), 1.2932605, -1.6092048, -1.5371683,
          ('慣犯', 0.035380281024169922), -1.6739159, 0.9763718, -0.8025273,
          ('詐欺罪', 0.290659717926065806), 0.2761859, 2.3008378, -0.7928667,
          ('信用調查', 0.51194528001768204834), 0.14855805, 3.2365885, -3.3151088,
          ('詐術', 0.05015530047607462197), -1.2938559, -2.1719165, 0.03638348,
          ('逃稅', 0.0308518656052148785), 1.6141753, -0.09968105, -2.897564,
          ('逃漏稅', 0.02632080166233987492), 2.1948044, 1.0531849, 1.6850697,
          ('賭博', 0.2827061774534225817), -0.9736342, 0.47620815, -0.53191006,
          2.4831314, -1.8135078, -0.09655156, -0.26880732, -1.5746521,
          3.447743, 0.7701042, -2.079731, -1.1631368, 0.29982182,
          -0.07716776, -3.0327997, 3.0748208, -1.4359124, 1.0085474,
          0.0964113, 0.32251447, 3.6297605, -1.2651255, -0.76287466,
          0.31288242, 0.9797535, 0.22678709, 0.11903057, 0.6259719,
          -2.007877, -0.33391798, -1.288079, 1.391751, 0.90182006,
          1.5843096, -0.94183683, -0.56292737, -1.1241164, -1.5128164,
          -1.3753456, -2.1157842, 1.3935559, -1.5951293, 1.7871724,
          -1.3674058, -0.77202106, 0.08307242, 0.7245565, 2.0208385,
```

```
-0.72172374, -1.9296253 , 0.6796764 , 2.4959662 , -0.91734695],  
In [25]: dtype=float32)
```

```
model.wv.most_similar("合約")
```

Out[25]:

```
[('合同', 0.7740179896354675),  
( '簽約', 0.6996151208877563),  
( '續約', 0.6664248704910278),  
( '租約', 0.616113007068634),  
( '簽下', 0.615209698677063),  
( '續簽', 0.6073216199874878),  
( '買斷', 0.5997638702392578),  
( '新東家', 0.5763907432556152),  
( '選擇權', 0.5660616159439087),  
( '勞動合同', 0.5596709251403809)]
```

## 計算 Cosine 相似度

In [26]:

```
model.wv.similarity("連結", "鏈接")
```

Out[26]:

0.709352

In [27]:

```
model.wv.similarity("連結", "陰天")
```

Out[27]:

-0.023800302

## 讀取模型

In [28]:

```
print(f"Loading {output_model}...")  
new_model = word2vec.Word2Vec.load(output_model)
```

Loading word2vec.zh.300.model...

In [29]:

```
model.wv.similarity("連結", "陰天") == new_model.wv.similarity("連結", "陰天")
```

Out[29]:

True

